

Pintos lab4

Youngjin Kwon

Contents

- Implementing file indexing APIs (FAT)
- Building hierarchical namespace
 - Hierarchical directory
- Modifying block access to use FAT
- Softlink

Indexing structure: FAT

- Formatting file system

```
void filesys_init (bool format) {  
    filesys_disk = disk_get (0, 1);  
    if (filesys_disk == NULL)  
        PANIC ("hd0:1 (hdb) not present");  
  
    inode_init ();  
  
#ifdef EFILESYS  
    fat_init ();  
  
    if (format)  
        do_format ();  
#endif  
}  
  
/* Formats the file system. */  
static void  
do_format (void) {  
    printf ("Formatting file system...");  
  
#ifdef EFILESYS  
    /* Create FAT and save it to the disk. */  
    fat_create ();  
    fat_close ();  
#else  
    free_map_create ();  
    if (!dir_create (ROOT_DIR_SECTOR, 16))  
        PANIC ("root directory creation failed");  
    free_map_close ();  
#endif  
  
    printf ("done.\n");  
}
```

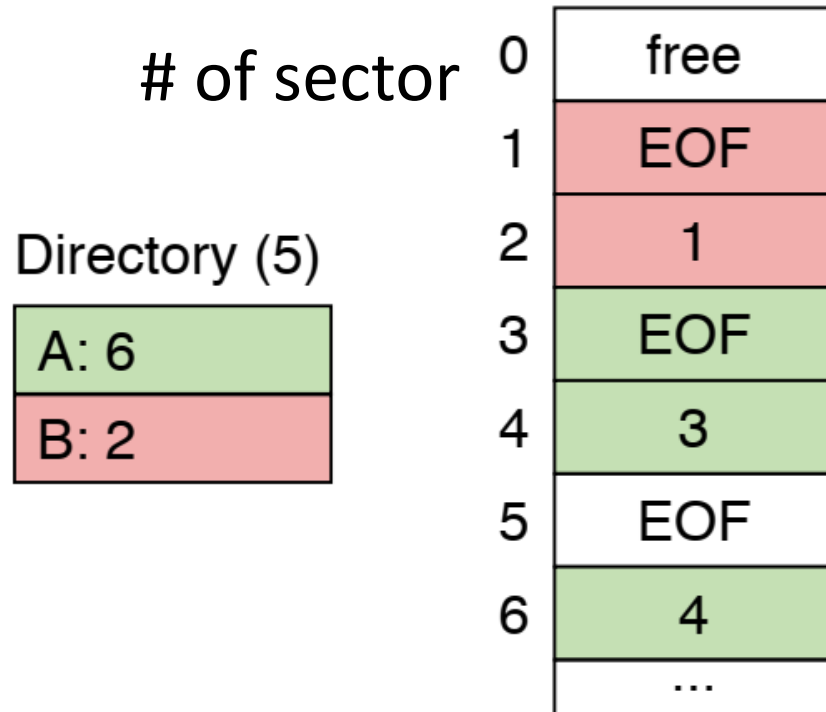
```
#define data_start(FATBS) (1 + ((FATBS).fat_sectors))  
#define data_sectors(FATBS) (((FATBS).total_sectors) - (data_start (FATBS)))
```

```
void fat_create (void) {  
    // Create FAT boot  
    fat_boot_create ();  
    fat_fs_init ();  
  
    // Create FAT table  
    fat_fs->fat = calloc (fat_fs->fat_length, sizeof (cluster_t));  
    if (fat_fs->fat == NULL)  
        PANIC ("FAT creation failed");  
  
    // Set up ROOT_DIR_CLST  
    fat_put (ROOT_DIR_CLUSTER, EOChain);  
  
    // Fill up ROOT_DIR_CLUSTER region with 0  
    uint8_t *buf = calloc (1, DISK_SECTOR_SIZE);  
    if (buf == NULL)  
        PANIC ("FAT create failed due to OOM");  
    disk_write (filesys_disk, cluster_to_sector (ROOT_DIR_CLUSTER), buf);  
    free (buf);  
}  
  
void fat_boot_create (void) {  
    unsigned int fat_sectors =  
        (disk_size (filesys_disk) - 1)  
        / (DISK_SECTOR_SIZE / sizeof (cluster_t) * SECTORS_PER_CLUSTER + 1) + 1;  
    fat_fs->bs = (struct fat_boot){  
        .magic = FAT_MAGIC,  
        .sectors_per_cluster = SECTORS_PER_CLUSTER,  
        .total_sectors = disk_size (filesys_disk),  
        .fat_start = 1,  
        .fat_sectors = fat_sectors,  
        root_dir_cluster = ROOT_DIR_CLUSTER,  
    };  
}
```


FAT review

- Size of fat table?


FAT (16-bit entries) • $\text{Fat_fs} \rightarrow \text{fat_length} * \text{sizeof}(\text{cluster_t})$



File A:



File B:



Implement APIs in `filesystem/fat.c` described in the project document

Hierarchical namespace (directory hierarchy)

- Current directory
 - Where to store the current directory? Struct thread. Use struct dir..
- You need to implement changing the current directory (e.g., the cd command)
 - Expand dir_create() to add . (current directory) and .. (parent directory) to each directory
 - Closely look at dir_add() or dir_remove() to figure out how directories are organized

Directory

- A directory is array of struct dir_entry

```
/* A single directory entry. */
struct dir_entry {
    disk_sector_t inode_sector;    /* Sector number of header. */
    char name[NAME_MAX + 1];      /* Null terminated file name. */
    bool in_use;                  /* In use or free? */
};
```

```
bool
filesys_create (const char *name, off_t initial_size) {
    disk_sector_t inode_sector = 0;
    struct dir *dir = dir_open_root ();
    bool success = (dir != NULL
                    && free_map_allocate (1, &inode_sector)
                    && inode_create (inode_sector, initial_size)
                    && dir_add (dir, name, inode_sector));
    if (!success && inode_sector != 0)
        free_map_release (inode_sector, 1);
    dir_close (dir);

    return success;
}
```

```
bool
dir_add (struct dir *dir, const char *name, disk_sector_t inode_sector) {
    struct dir_entry e;
    off_t ofs;
    bool success = false;

    ASSERT (dir != NULL);
    ASSERT (name != NULL);

    /* Check NAME for validity. */
    if (*name == '\0' || strlen (name) > NAME_MAX)
        return false;

    /* Check that NAME is not in use. */
    if (lookup (dir, name, NULL, NULL))
        goto done;

    /* Set OFS to offset of free slot.
     * If there are no free slots, then it will be set to the
     * current end-of-file.

     * inode_read_at() will only return a short read at end of file.
     * Otherwise, we'd need to verify that we didn't get a short
     * read due to something intermittent such as low memory. */
    for (ofs = 0; inode_read_at (dir->inode, &e, sizeof e, ofs) == sizeof e;
         ofs += sizeof e)
        if (!e.in_use)
            break;

    /* Write slot. */
    e.in_use = true;
    strcpy (e.name, name);
    e.inode_sector = inode_sector;
    success = inode_write_at (dir->inode, &e, sizeof e, ofs) == sizeof e;

done:
    return success;
}
```

Write data to disk using FAT

Handle the case where directory and file grows

```
off_t
inode_write_at (struct inode *inode, const void *buffer_, off_t size,
               off_t offset) {
    const uint8_t *buffer = buffer_;
    off_t bytes_written = 0;
    uint8_t *bounce = NULL;

    if (inode->deny_write_cnt)
        return 0;

    while (size > 0) {
        /* Sector to write, starting byte offset within sector. */
        disk_sector_t sector_idx = byte_to_sector (inode, offset);
        int sector_ofs = offset % DISK_SECTOR_SIZE;

        /* Bytes left in inode, bytes left in sector, lesser of the two. */
        off_t inode_left = inode_length (inode) - offset;
        int sector_left = DISK_SECTOR_SIZE - sector_ofs;
        int min_left = inode_left < sector_left ? inode_left : sector_left;

        /* Number of bytes to actually write into this sector. */
        int chunk_size = size < min_left ? size : min_left;
        if (chunk_size <= 0)
            break;

        if (sector_ofs == 0 && chunk_size == DISK_SECTOR_SIZE) {
            /* Write full sector directly to disk. */
            disk_write (filesystem_disk, sector_idx, buffer + bytes_written);
        } else {
            /* We need a bounce buffer. */
            if (bounce == NULL) {
                bounce = malloc (DISK_SECTOR_SIZE);
                if (bounce == NULL)
                    break;
            }

            /* If the sector contains data before or after the chunk
               we're writing, then we need to read in the sector
               first. Otherwise we start with a sector of all zeros. */
            if (sector_ofs > 0 || chunk_size < sector_left)
                disk_read (filesystem_disk, sector_idx, bounce);
            else
                memset (bounce, 0, DISK_SECTOR_SIZE);
            memcpy (bounce + sector_ofs, buffer + bytes_written, chunk_size);
            disk_write (filesystem_disk, sector_idx, bounce);
        }
    }
}
```

```
static disk_sector_t
byte_to_sector (const struct inode *inode, off_t pos) {
    ASSERT (inode != NULL);
    if (pos < inode->data.length)
        return inode->data.start + pos / DISK_SECTOR_SIZE;
    else
        return -1;
}
```

Update this part to use
FAT you implement

Also, modify `inode_create()` which
creates new inode sector and
`inode_open()` which read inode sectors
to use FAT

Open a file (or directory) with absolute path

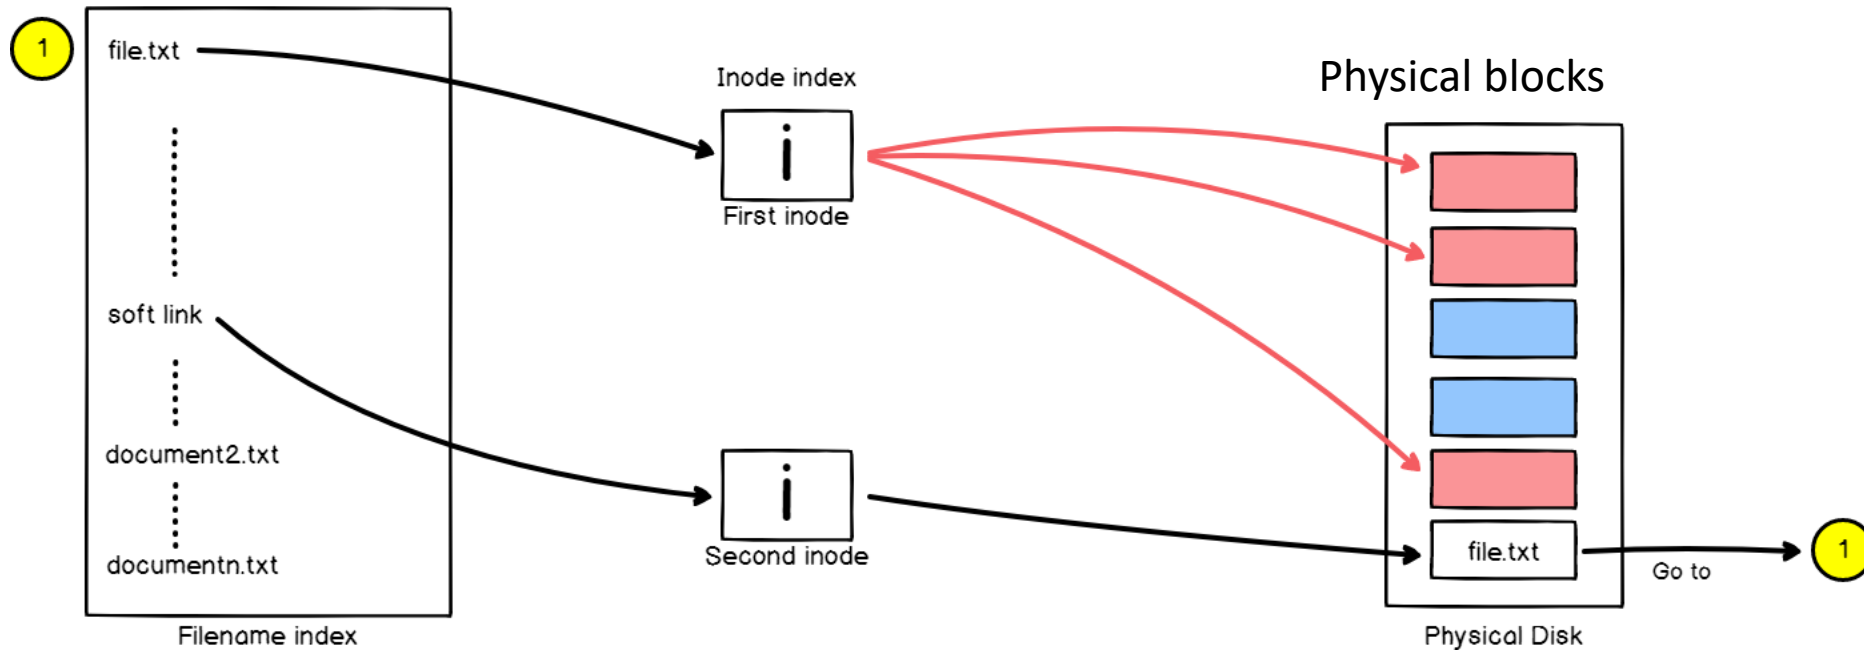
- Open("/a/b/c/d") or remove("/a/b/c/d")
 - File system should walk directories from "/" to "/a/b/c/" to find d

```
struct file *  
filesystem_open (const char *name) {  
    struct dir *dir = dir_open_root ();  
    struct inode *inode = NULL;  
  
    if (dir != NULL)  
        dir_lookup (dir, name, &inode);  
    dir_close (dir);  
  
    return file_open (inode);  
}
```

- Extend filesystem_open() and filesystem_remove() to walk a directory hierarchy given by *name

Softlink

Understanding Soft Links



Extend inode to support soft link

- Symlink creates a new inode
 - Contents are duplicated from the target inode
 - Treat symlink inode as a normal inode
 - E.g., symlink name is in a directory, and inode # points to a symlink inode
 - Deletion of symlink does not delete file
 - Symlink can point to a directory
 - Symlink remains even if a target file is deleted (dangling symlink)

```
devconnected@debian-10:~$ ls -l shortcut
lrwxrwxrwx 1 devconnected devconnected 8 Aug 13 16:18 shortcut -> file.txt
devconnected@debian-10:~$ rm file.txt
devconnected@debian-10:~$ ls -l shortcut
lrwxrwxrwx 1 devconnected devconnected 8 Aug 13 16:18 shortcut -> file.txt
```